

# Data Science

Marc Tommasi

September 26, 2023

# Outline

## 1 Linear Regression

# Outline

## 1 Linear Regression

# Main ideas

- Used to explain as well as to predict. Example: explain fuel consumption as a function of car weight.
- Take the class of straight lines  $y = ax + b$ . Find  $a$  and  $b$  given weights ( $x$ ) and observed consumption ( $y$ ).
- We also call  $b$  *intercept* (ordonnée à l'origine) and  $a$  is a coefficient (which defines the slope of the line).
- If we know that the data  $x$  are centered at 0, then the intercept is zero.
- Generalization to the case where data are represented in a space of dimension  $d$ . Example: predict fuel consumption as a function of car weight and engine power, weather, speed, etc.
- We take a class  $\mathcal{H}$  of linear functions: we look for  $w$ ; such that  $y = w_0 + w_1x_1 + \dots + w_dx_d$ .
- Note: we can pose  $x_0 = 1$  and compute

$$y = w_0x_0 + w_1x_1 + \dots + w_dx_d = \mathbf{w}^\top \mathbf{x} = \langle \mathbf{w}, \mathbf{x} \rangle$$

# Formal setting

- $\mathcal{X} \subseteq \mathbb{R}^d$ ,  $\mathcal{Y} \subseteq \mathbb{R}$
- $\mathcal{H} = \{\mathbf{x} \rightarrow \langle \mathbf{w}, \mathbf{x} \rangle \mid \mathbf{w} \in \mathbb{R}^d\}$  is the class of hypothesis.
- There are  $m$  examples of the form  $((x_1, x_2, \dots, x_d), y) = (\mathbf{x}, y)$ :

$$\mathcal{S} = \{(\mathbf{x}^{(1)}, y_1), (\mathbf{x}^{(2)}, y_2), \dots, (\mathbf{x}^{(m)}, y_m)\}$$

- $\mathcal{S}$  can be represented by a matrix  $X$  where each row is an instance of  $\mathbf{x}$
- We use the quadratic loss (squared error) to measure the empirical error:

$$\ell(h, (\mathbf{x}, y)) = (h(\mathbf{x}) - y)^2$$

# ERM and the objective function

- The aim is to minimize the squared error for each example, and the loss function, the **objective function**, is the Mean squared error (MSE)

$$\frac{1}{m} \sum_i (\langle \mathbf{w}, \mathbf{x}^{(i)} \rangle - y_i)^2 = \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

- The **ERM principle** is to find the parameters  $\mathbf{w}$  that minimize the objective function on  $S$ :

$$\operatorname{argmin}_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

- The objective is convex: there is a global minimum and it can be calculated analytically.

# Derivatives and gradient

- The derivative  $f'$  of a function  $f : \mathbb{R} \rightarrow \mathbb{R}$  evaluated at  $x$  is defined by

$$f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- The derivative cancels for an  $x$  that minimizes  $f$ .
- For  $f : \mathbb{R}^d \rightarrow \mathbb{R}$ , we can consider the derivative with respect to each component of  $x = (x_1, \dots, x_d)$ . This is denoted  $\frac{\partial f}{\partial x_i}$ .
- The  $d$ -dimensional vector of all these derivative functions is the gradient  $\nabla f = \left( \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_d} \right)$ .
- Remember: the derivative of  $f(x) = x^n$  is  $f'(x) = nx^{n-1}$  and the derivative of a constant is zero;

# Conclusion

- We can find an analytical solution which will be  $(X^T X)^{-1} X^T y$ .
- Informally, we're looking for  $Xw = y$ , which we can rewrite as  $X^T X w = X^T y$ , then  $w = (X^T X)^{-1} X^T y$ .
- Note
  - ▶ The matrix  $(X^T X)$  may not be invertible, so we can instead calculate  $X^\dagger y$  where  $X^\dagger$  is the Moore-Penrose pseudo-inverse, which is calculated by an SVD (If  $X = U \Sigma V^T$  then  $X^\dagger = V \Sigma^\dagger U^T$  and  $\Sigma^\dagger$  is calculated by taking the inverse of the values except for the 0s, which remain 0).

# Complexity

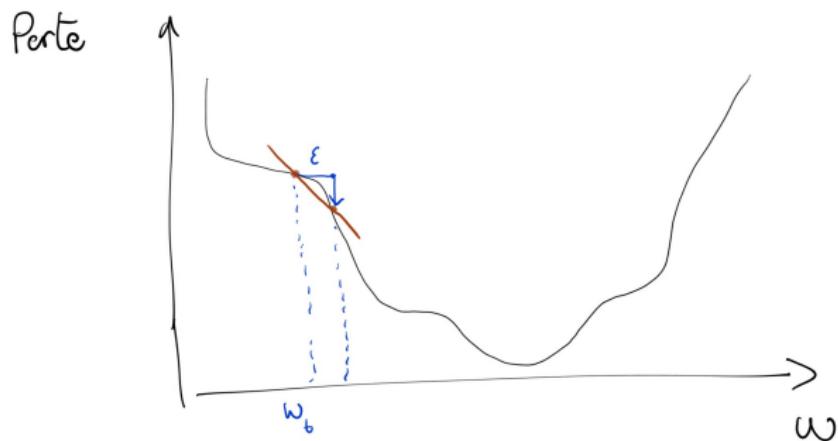
- The matrix inversion is costly for very large datasets
- Matrix multiplication is  $O(n^3)$  for simple algorithms ( $O(n^{2.4})$  for the best ones).
- The gradient descent approach yields a quadratic algorithm that works quite well in practice.
- (and can be accelerated with GPUs...)

# Gradient descent I

- The aim is to reduce the empirical error step by step: find the parameter modification that will make the loss function take on a smaller value.
- Idea of descending through the steepest slope on the curve of loss as a function of parameters  $w$

# Gradient descent II

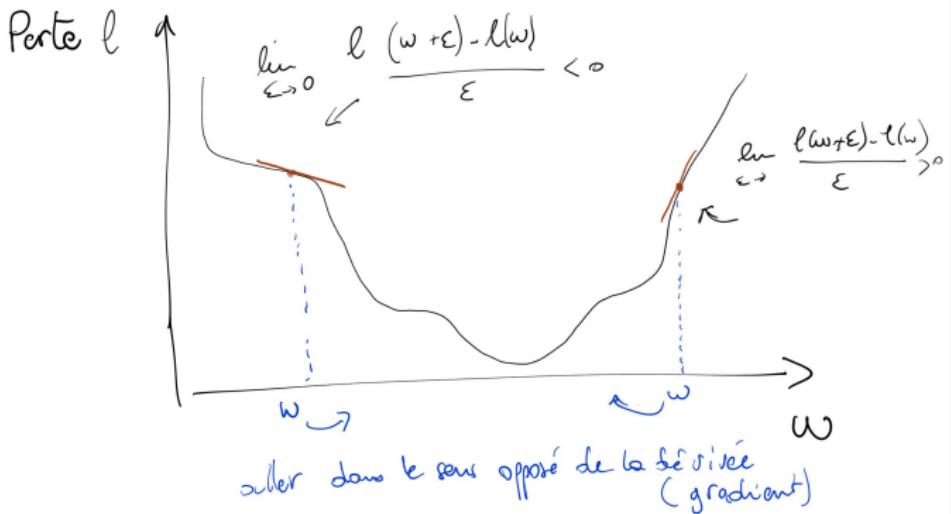
## Derivative direction



Quick reminder about derivatives

# Gradient descent III

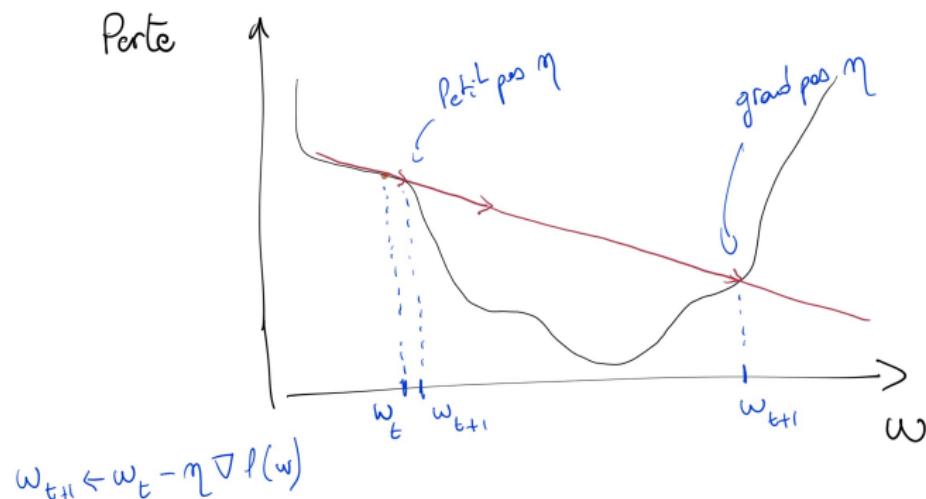
## Direction



The gradient gives the direction: opposite to the gradient

# Gradient descent IV

## Step by step



We're taking a step in this direction, by a factor yet to be specified: the learning rate.

# Algorithm

- Start with random weights, which define a point on the curve of the objective function.
- Calculate the derivative (gradient) of the objective function at this point
- One step is taken in the opposite direction to this gradient. The length of this step is determined by the learning rate.
- A termination check is performed on the norm of the gradient (tolerance) or on the error.
- Notes
  - ▶ In the convex case (as with MSE), a global minimum is guaranteed.
  - ▶ It is preferable to have normalized dimensions (use StandardScaler for example).
  - ▶ Dynamic gradient step can be used: the learning rate value decreases slowly (on the order of  $1/t$ ).

# More formally

- Input: Sample  $S = \{(\mathbf{x}^{(1)}, y_1), (\mathbf{x}^{(2)}, y_2), \dots, (\mathbf{x}^{(m)}, y_m)\}$ , a loss function  $f$ , a learning rate  $\eta$
- Init:  $\mathbf{w}^{(0)}$  receives random values
- Repeat until convergence
  - ▶  $t = t + 1$
  - ▶  $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla_{\mathbf{w}^{(t)}} f(\mathbf{w}, S)$
- Output :  $\mathbf{w}^{(t)}$

# Batch and stochastic

## Gradient descent (Batch Gradient Descent)

- The algorithm requires each iteration to pass over the entire data set (reminder: the loss function is a sum over all examples).
- The cost of this calculation can still be prohibitive
- Sometimes difficult to escape from local minima

## Stochastic Gradient Descent (SGD)

- A single example is selected at each round
- Advantages
  - ▶ fast
  - ▶ if the loss function is very unstable, SGD can get out of local minima

## Mini Batch Gradient Descent

- a subset of the examples is selected
- suitable for GPUs, as we can now perform small matrix operations on a batch of data

# Polynomial regression

- Linear models have been constructed (straight lines, hyperplanes, etc.).
- This class of hypotheses may be ill-suited to non-linear data.
- We can get around this problem by adding new attributes to the data: square, cube, pairwise products of attributes, etc.
- In the 1-dimension case, let's solve

$$y = w_0 + w_1 x + w_2 x^2 + \cdots + w_k x^k$$

- The data are transformed, but the same linear regression algorithm is applied!
- The same applies if  $x$  is a vector in  $\mathbb{R}^d$  and products can be made between different attributes ( $x_i x_j$ ,  $x_i x_j x_k \dots$ ).
- NB: we can avoid materializing the construction of all these new attributes using **kernel** techniques (used in SVMs, ...).

# Technically

- with PolynomialFeatures.
- we can use Pipeline class objects to compose different stages
  - ▶ feature transformation

```
estimators = [('reduce_dim', PCA()), ('clf', SVC())]
pipe = Pipeline(estimators)
# two _ to access params
pipe.set_params(clf__C=10)
```
  - ▶ application of a model with pipe.fit, (score, predict,... )