

# Data Science

Marc Tommasi

October 10, 2023

# Outline

## 1 Linear Regression

# Outline

## 1 Linear Regression

# Main ideas

- Used to explain as well as to predict. Example: explain fuel consumption as a function of car weight.
- Take the class of straight lines  $y = ax + b$ . Find  $a$  and  $b$  given weights ( $x$ ) and observed consumption ( $y$ ).
- We also call  $b$  *intercept* (ordonnée à l'origine) and  $a$  is a coefficient (which defines the slope of the line).
- If we know that the data  $x$  are centered at 0, then the intercept is zero.
- Generalization to the case where data are represented in a space of dimension  $d$ . Example: predict fuel consumption as a function of car weight and engine power, weather, speed, etc.
- We take a class  $\mathcal{H}$  of linear functions: we look for  $w$ ; such that  $y = w_0 + w_1x_1 + \dots + w_dx_d$ .
- Note: we can pose  $x_0 = 1$  and compute

$$y = w_0x_0 + w_1x_1 + \dots + w_dx_d = \mathbf{w}^\top \mathbf{x} = \langle \mathbf{w}, \mathbf{x} \rangle$$

# Formal setting

- $\mathcal{X} \subseteq \mathbb{R}^d$ ,  $\mathcal{Y} \subseteq \mathbb{R}$
- $\mathcal{H} = \{\mathbf{x} \rightarrow \langle \mathbf{w}, \mathbf{x} \rangle \mid \mathbf{w} \in \mathbb{R}^d\}$  is the class of hypothesis.
- There are  $m$  examples of the form  $((x_1, x_2, \dots, x_d), y) = (\mathbf{x}, y)$ :

$$\mathcal{S} = \{(\mathbf{x}^{(1)}, y_1), (\mathbf{x}^{(2)}, y_2), \dots, (\mathbf{x}^{(m)}, y_m)\}$$

- $\mathcal{S}$  can be represented by a matrix  $X$  where each row is an instance of  $\mathbf{x}$
- We use the quadratic loss (squared error) to measure the empirical error:

$$\ell(h, (\mathbf{x}, y)) = (h(\mathbf{x}) - y)^2$$

# ERM and the objective function

- The aim is to minimize the squared error for each example, and the loss function, the **objective function**, is the Mean squared error (MSE)

$$\frac{1}{m} \sum_i (\langle \mathbf{w}, \mathbf{x}^{(i)} \rangle - y_i)^2 = \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

- The **ERM principle** is to find the parameters  $\mathbf{w}$  that minimize the objective function on  $S$ :

$$\operatorname{argmin}_{\mathbf{w}} \frac{1}{m} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

- The objective is convex: there is a global minimum and it can be calculated analytically.

# Derivatives and gradient

- The derivative  $f'$  of a function  $f : \mathbb{R} \rightarrow \mathbb{R}$  evaluated at  $x$  is defined by

$$f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon}$$

- The derivative cancels for an  $x$  that minimizes  $f$ .
- For  $f : \mathbb{R}^d \rightarrow \mathbb{R}$ , we can consider the derivative with respect to each component of  $x = (x_1, \dots, x_d)$ . This is denoted  $\frac{\partial f}{\partial x_i}$ .
- The  $d$ -dimensional vector of all these derivative functions is the gradient  $\nabla f = (\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_d})$ .
- Remember: the derivative of  $f(x) = x^n$  is  $f'(x) = nx^{n-1}$  and the derivative of a constant is zero;

# Conclusion

- We can find an analytical solution which will be  $(X^T X)^{-1} X^T y$ .
- Informally, we're looking for  $Xw = y$ , which we can rewrite as  $X^T X w = X^T y$ , then  $w = (X^T X)^{-1} X^T y$ .
- Note
  - ▶ The matrix  $(X^T X)$  may not be invertible, so we can instead calculate  $X^\dagger y$  where  $X^\dagger$  is the Moore-Penrose pseudo-inverse, which is calculated by an SVD (If  $X = U \Sigma V^T$  then  $X^\dagger = V \Sigma^\dagger U^T$  and  $\Sigma^\dagger$  is calculated by taking the inverse of the values except for the 0s, which remain 0).

# Complexity


- The matrix inversion is costly for very large datasets
- Matrix multiplication is  $O(n^3)$  for simple algorithms ( $O(n^{2.4})$  for the best ones).
- The gradient descent approach yields a quadratic algorithm that works quite well in practice.
- (and can be accelerated with GPUs...)

# Gradient descent I

- The aim is to reduce the empirical error step by step: find the parameter modification that will make the loss function take on a smaller value.
- Idea of descending through the steepest slope on the curve of loss as a function of parameters  $w$

# Gradient descent II


## Derivative direction



Quick reminder about derivatives

# Gradient descent III

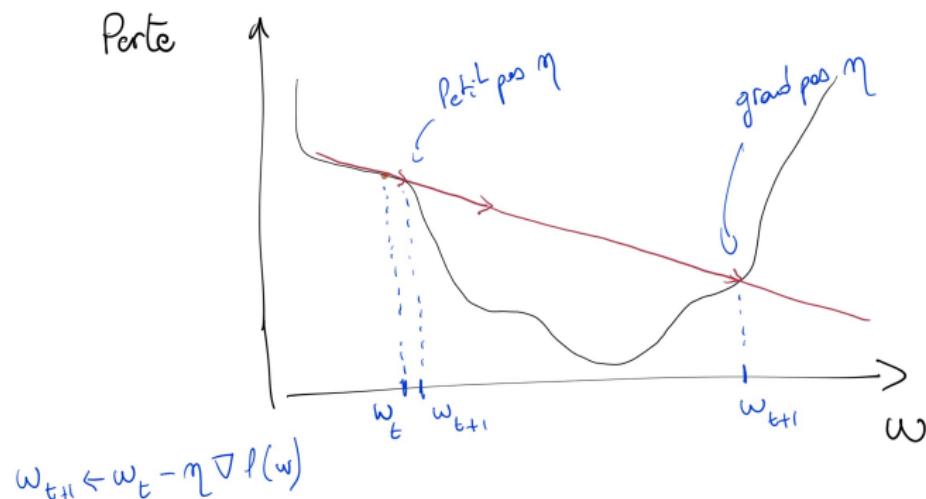
## Direction



The gradient gives the direction: opposite to the gradient

# Gradient descent IV

## Step by step



We're taking a step in this direction, by a factor yet to be specified: the learning rate.

# Algorithm

- Start with random weights, which define a point on the curve of the objective function.
- Calculate the derivative (gradient) of the objective function at this point
- One step is taken in the opposite direction to this gradient. The length of this step is determined by the learning rate.
- A termination check is performed on the norm of the gradient (tolerance) or on the error.
- Notes
  - ▶ In the convex case (as with MSE), a global minimum is guaranteed.
  - ▶ It is preferable to have normalized dimensions (use StandardScaler for example).
  - ▶ Dynamic gradient step can be used: the learning rate value decreases slowly (on the order of  $1/t$ ).

# More formally

- Input: Sample  $S = \{(\mathbf{x}^{(1)}, y_1), (\mathbf{x}^{(2)}, y_2), \dots, (\mathbf{x}^{(m)}, y_m)\}$ , a loss function  $f$ , a learning rate  $\eta$
- Init:  $\mathbf{w}^{(0)}$  receives random values
- Repeat until convergence
  - ▶  $t = t + 1$
  - ▶  $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla_{\mathbf{w}^{(t-1)}} f(\mathbf{w}^{(t-1)}, S)$
- Output :  $\mathbf{w}^{(t)}$

# Batch and stochastic

## Gradient descent (Batch Gradient Descent)

- The algorithm requires each iteration to pass over the entire data set (reminder: the loss function is a sum over all examples).
- The cost of this calculation can still be prohibitive
- Sometimes difficult to escape from local minima

## Stochastic Gradient Descent (SGD)

- A single example is selected at each round
- Advantages
  - ▶ fast
  - ▶ if the loss function is very unstable, SGD can get out of local minima

## Mini Batch Gradient Descent

- a subset of the examples is selected
- suitable for GPUs, as we can now perform small matrix operations on a batch of data

# Polynomial regression

- Linear models have been constructed (straight lines, hyperplanes, etc.).
- This class of hypotheses may be ill-suited to non-linear data.
- We can get around this problem by adding new attributes to the data: square, cube, pairwise products of attributes, etc.
- In the 1-dimension case, let's solve

$$y = w_0 + w_1 x + w_2 x^2 + \cdots + w_k x^k$$

- The data are transformed, but the same linear regression algorithm is applied!
- The same applies if  $x$  is a vector in  $\mathbb{R}^d$  and products can be made between different attributes ( $x_i x_j$ ,  $x_i x_j x_k \dots$ ).
- NB: we can avoid materializing the construction of all these new attributes using **kernel** techniques (used in SVMs, ...).

# Technically

- with PolynomialFeatures.
- we can use Pipeline class objects to compose different stages
  - ▶ feature transformation

```
estimators = [('reduce_dim', PCA()), ('clf', SVC())]
pipe = Pipeline(estimators)
# two _ to access params
pipe.set_params(clf__C=10)
```
  - ▶ application of a model with pipe.fit, (score, predict,... )

# Error decomposition

## Reminder: Complexity/variance bias

- bias: assumption about the function class, for example. Subject to underfitting if bias is too strong
- variance/complexity: variation in function class due to sensitivity to data. Subject to overfitting if the variance is too high.
- NB: there always remains an irreducible error due to data quality in the error decomposition.

## Error decomposition

- $L_D(h_S)$  is the sum of the approximation error + the estimation error
- Approximation error
  - ▶ Does not depend on  $S$  but on the choice of  $\mathcal{H}$ .
  - ▶ Decreases with increasing complexity of  $\mathcal{H}$ .
- Estimation error
  - ▶ decreases if we have more examples in  $S$  (we reduce the variance...)

# Illustration of polynomial regression

- the technique of adding numerous attributes increases complexity and variance.
- the degree of the  $p$  polynomials,  $x_i, x_i^2, \dots x_i^p$ ; the sizes of the products  $x_i x_j, \dots$  are **hyper-parameters** to be adjusted to control bias/complexity.
- we must try to control this:
  - ▶ tracing **learning curves**,
  - ▶ **regularization**

# Methodology

- Hyper-parameters must be set on the learning sample.


## Remember

- The test set should **never** be looked at, except to evaluate the true/generalization error.

## Model selection techniques

- Training set split into learning and validation or cross-validation
- cross-validation,
- leave-one-out,...
- [https://scikit-learn.org/stable/modules/classes.html#module-sklearn.model\\_selection](https://scikit-learn.org/stable/modules/classes.html#module-sklearn.model_selection)

# Cross-validation



# In practice

## Learning curves

- calculate the error in training and validation according to sample size or model complexity.
- examining the curve can help select the right parameters



## Grid Search

- vary the set of hyper-parameters and perform an exhaustive search for the best set of parameters

# Regularization I

## Ridge Regression

- Tikonov regularization: adds the square of the norm of the parameter vector ( $+ \alpha \frac{1}{2} \sum_i w_i^2$ ). This is also known as a  $\ell_2$  penalty.
- It's best to normalize the data before applying this regularization.
- Increases the bias and reduces the variance
- A closed form can be calculated, but SGD can also be applied.
- In sklearn, class Ridge or SGDRegressor with penalty which is "l2" with parameter alpha.

# Regularization II

## Lasso

- This time it's the  $\ell_1$  standard, the sum of absolute values ( $\alpha \sum_i |\mathbf{w}_i|$ ).
- Removes the least relevant attributes: a sparse model is obtained by variable selection.
- The function is not derivable in 0 (right derivative different from left derivative).
- We can use a sub-differential (sub-gradient, subgradient) around 0, taking the sign if  $\mathbf{w}$  is non-zero and 0 otherwise.
- Class Lasso with parameter alpha.

## Elastic net

- this is the convex sum of the two regularizations (a ratio is applied between the 2).
- Class ElasticNet with parameters ratio and alpha.

# Regularization III

## Early stopping

- Limits the number of iterations. This is a form of “algorithmic” regularization.
- We stop when the validation error rises.
- Not easy to identify, as the curves are not smooth!
- parameter `earlystopping` in SGD-type algos.
- can also be simulated with the `warm_start` parameter, which allows optimization to continue where it left off (with `max_iter` to set the number of iterations).
- `SGDRegressor` with zero regularization takes `penalty=None`.

# Regression and classification

- Classification problem :
  - ▶ predict a value in a finite set  $\mathcal{Y}$ .
  - ▶ in general, the cardinality of  $\mathcal{Y}$  is small
  - ▶ values in  $\mathcal{Y}$  are not assumed to be ordered
- Many methods use **binary classification** ( $\mathcal{Y} = \{-1, +1\}$ ).
- If the cardinality of  $\mathcal{Y}$  is greater than 2: **multi-class problem**.
- Many methods can be used in regression or classification, with a few variations.
  - ▶ Example: predict using  $\text{sign}(\langle w, x \rangle)$ .


# Logistic regression

- Logistic regression is not regression but classification.
- The aim is to estimate the probability  $P$  of having a class (e.g.  $+1, -1$ ) given a data description  $x$ .
- We will learn a function  $h_w$  parametrized by a set of parameters  $w$ .
- The value of  $h_w(x) \in [0, 1]$  and  $\int_x h_w(x) dx = 1$  (a little abuse of notation)
- Hence,  $h_w$  can be used to compute  $P(Y = 1 | x)$  as  $h_w(X = x)$

# Logistic function

- We use the sigmoidal logistic function:

$$h_w(x) = \sigma(w^\top x) \text{ with } \sigma(x) = \frac{1}{1 + \exp(-x)}$$



- It's a continuous function that can be derived between 0 and 1 and passes through 0.5 in 0.
- The decision threshold is 1/2.

# Identity

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})}$$

$$\begin{aligned}1 - h_{\mathbf{w}}(\mathbf{x}) &= 1 - \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} \\&= \frac{\exp(-\mathbf{w}^\top \mathbf{x})}{1 + \exp(-\mathbf{w}^\top \mathbf{x})} \\&= \frac{1}{1 + \frac{1}{\exp(-\mathbf{w}^\top \mathbf{x})}} \\&= \frac{1}{1 + \exp(\mathbf{w}^\top \mathbf{x})}\end{aligned}$$

Remarque :  $\log(h_{\mathbf{w}}(\mathbf{x})) = -\log(1 + \exp(-\mathbf{w}^\top \mathbf{x}))$

## Training: explanation by likelihood

- We have the estimate of the probability of having  $y = +1$  which is  $h_w(\mathbf{x})$  and  $y = -1$  which is  $1 - h_w(\mathbf{x})$ .
- The idea is to find the  $w$  that will maximize the likelihood of observing our sample.
- So for each pair  $(\mathbf{x}, y)$ , we maximize the expression ( $\mathbb{1}_t$  is 1 if  $t$  is true and 0 otherwise).

$$h_w(\mathbf{x})^{\mathbb{1}_{y=+1}}(1 - h_w(\mathbf{x}))^{\mathbb{1}_{y=-1}}$$

- We often take the log-likelihood, for ease of calculation

$$\mathbb{1}_{y=+1} \log(h_w(\mathbf{x})) + \mathbb{1}_{y=-1} \log(1 - h_w(\mathbf{x}))$$

- This is equivalent to minimizing the opposite, known as the loss:

$$\begin{aligned}\ell(h_w, (\mathbf{x}, y)) &= -\mathbb{1}_{y=+1} \log(h_w(\mathbf{x})) - \mathbb{1}_{y=-1} \log(1 - h_w(\mathbf{x})) \\ &= \mathbb{1}_{y=+1} \log(1 + \exp(-\mathbf{w}^\top \mathbf{x})) + \mathbb{1}_{y=-1} \log(1 + \exp(\mathbf{w}^\top \mathbf{x})) \\ &= \log(1 + \exp(-y \mathbf{w}^\top \mathbf{x}))\end{aligned}$$

## Training: intuitive shortcut

- We want  $h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1+\exp(-\mathbf{w}^\top \mathbf{x})}$  to be large when  $y = 1$  (and  $1 - h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1+\exp(\mathbf{w}^\top \mathbf{x})}$  large if  $y = -1$ ).
- We therefore want a loss to grow with  $\frac{1}{1+\exp(y\mathbf{w}^\top \mathbf{x})}$  or  $1 + \exp(-y\mathbf{w}^\top \mathbf{x})$ .
- the associated loss function is the log of this quantity

$$\ell(h_{\mathbf{w}}, (\mathbf{x}, y)) = \log(1 + \exp(-y\mathbf{w}^\top \mathbf{x}))$$

- We apply the ERM principle to find  $\mathbf{w}$ .
- It's convex! but there's no closed form, but we use gradient descent to find the minimum.
- the same types of regularization can be applied as in linear regression
- This is the LogisticRegression class (the regularization parameter is  $C$ , a large value of  $C$  results in weak regularization).

# Multi-class problems

- one set of parameters per class:  $s_k(\mathbf{x}) = (\mathbf{w}^{(k)})^\top \mathbf{x}$  gives a score per class.
- we want probabilities therefore between 0 and 1 and the sum over all classes is 1 :

$$p_k = \frac{\exp(s_k(\mathbf{x}))}{\sum_j \exp(s_j(\mathbf{x}))}$$

- the largest value is strongly reinforced by the exponential: a continuous version of the calculation of a max, the **softmax**.
- we then predict the  $k$  that maximizes this softmax.
- in the binary case, we find  $h_{\mathbf{w}}$  by posing  $\mathbf{w} = \mathbf{w}^{(1)} - \mathbf{w}^{(2)}$  (exercise).
- class LogisticRegression with parameter `multi_class`.