

# Data Science

Marc Tommasi

November 7, 2023

# Outline

1 Decision Trees

2 Evaluation of classification models

3 Ensemble methods

# Outline

1 Decision Trees

2 Evaluation of classification models

3 Ensemble methods

# Principles

- Classifier: a function from  $\mathcal{X}$  to  $\mathcal{Y}$
- Model in the form of a set of successive decision rules, represented in a tree
- Model easy to interpret when the tree is small
- Example on iris dataset, 3 classes: setosa, virginica, versicolor; 150 examples; attributes sepal length, sepal width, petal length, petal width.

# Example



# Inference

- Start from the root, with an example (e.g. 3, 2, 3, 2)
- We go through the tests, each test cutting the space into 2 parts.
- We thus designate a recursive partitioning of the description space
- Each leaf gives a label to a part.
- Good visualization in [Jake VandenPlas's book](#).

# Algorithms

- This method introduces two biases: choice of function class + algorithmic bias.
- The algorithmic bias stems from a greedy heuristic:
  - ▶ the tree is built from root to leaves,
  - ▶ at each stage, a node corresponding to part of the data is developed
  - ▶ the best test is selected according to a gain criterion
  - ▶ this choice is no longer questioned
- Several algorithms exist: ID3, C4.5, C5, CART, etc.
- Sklearn implements CART

# Explanations with ID3

- We consider binary attributes  $A = \{0, 1, \dots, p\}$

```
def id3(S,A):  
    """ S: dataset, A: set of attributes """  
    if all examples in S have the same class or A is empty:  
        retrun a leaf labelled with the majority class
```

Let  $j$  be the attribute that maximizes the gain

```
t_l = id3({(x,y) in S st x_j=0}, A\{j})  
t_r = id3({(x,y) in S st x_j=1}, A\{j})  
return the tree whose root node is the test x_j=1  
    and the left children is t_l (case False)  
and the right children is t_r (case True).
```

# Compute the gain

- **Gain:** difference observed between absence and presence of the test in the tree.
- Notation: Let  $\mathbb{P}_S[F]$  be the probability of  $F$  when samples are drawn iid in  $S$ 
  - ▶ (if a node with test  $x_i$  corresponds to  $m$  examples and  $m_l$  go to the left,  $m_r$  go to the right, then  $\mathbb{P}_S[x_i = 1] = m_r/m$ )
- the gain is computed with a function  $C$  to be defined:

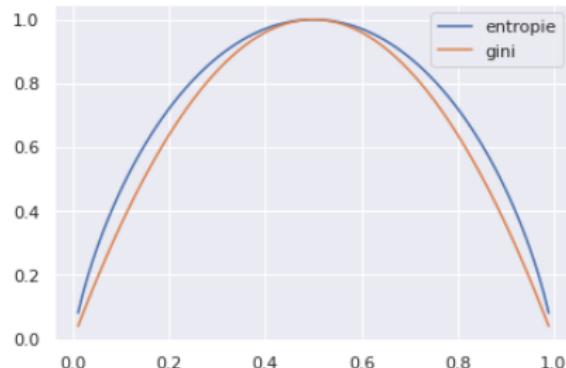
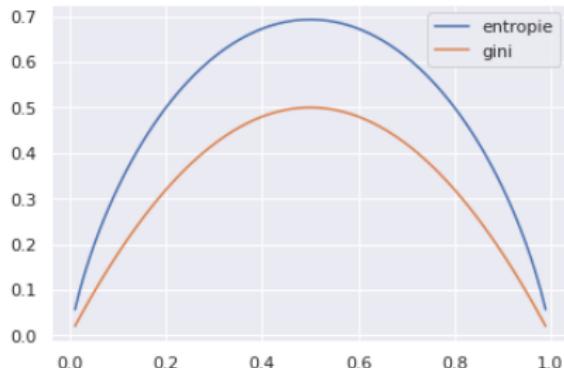
$$\text{Gain}(S, i) = C(\mathbb{P}_S[y = 1]) - (\mathbb{P}_S[x_i = 1]C(\mathbb{P}_S[y = 1 | x_i = 1]) + \mathbb{P}_S[x_i = 0]C(\mathbb{P}_S[y = 1 | x_i = 0]))$$

$C(a)$  can be:

- the training error:  $C(a) = \min(a, 1 - a)$
- the information gain (entropy):  $C(a) = -a \log(a) - (1 - a) \log(1 - a)$
- the Gini score:  $C(a) = 2a(1 - a)$

# Entropy and Gini

- Will favor tests that achieve the best separation, ie. when  $\mathbb{P}_S[y | x_i]$  is close to 0 or 1.



- $-a \log(a) - (1-a) \log(1-a)$   
et  $2a(1-a)$
- $-a \log_2(a) - (1-a) \log_2(1-a)$   
et  $4a(1-a)$

# Limiting overfitting

- Bias/complexity trade-off: the greater the depth, the more complex the function class, the lower the bias but the greater the tendency of decision trees to overfitting.
- depth bound (as in sklearn, parameter `max_depth`)
- pruning consists of removing branches: replacing a node with a class label.
  - ▶ Bottom-up approach with a statistical test ( $\xi^2$  or error evaluation).

# The case of continuous attributes

- discretization considers all possible thresholds observed on the training sample
- the calculation for these  $m$  possible tests could be  $O(dm^2)$  but can be reduced to  $O(dm \log(m))$ .

# Regression Trees

- the examples arriving in a leaf are averaged to determine the value to be predicted
- the cost function used to build the tree is, for example, the MSE

# Advantages and disadvantages

- model readability
- In theory, high algorithmic complexity to find the best tree, but fast heuristic approach in practice.
- difficulty in setting the max depth or other criteria to avoid overfitting

# Outline

1 Decision Trees

2 Evaluation of classification models

3 Ensemble methods

# Main measures

- **Taux d'erreurs** (accuracy) : the ratio of the number of wrongly classified examples and the total number of examples.
- **Confusion Matrix** : number of examples according to prediction and true class.

	Positifs	Negatifs
Prediction Positive	TP	FP
Prediction Negative	FN	TN

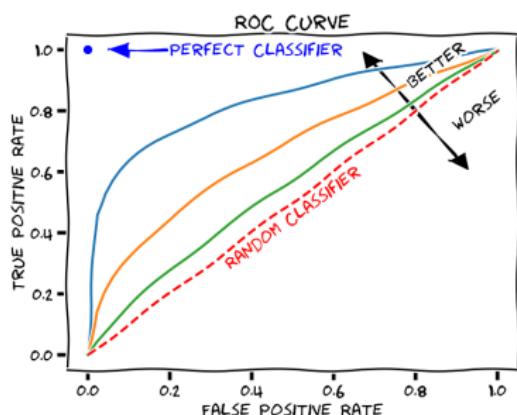
- **Precision** ( $\frac{TP}{TP+FP}$ ) and **Recall** ( $\frac{TP}{TP+FN}$ )
- The harmonic mean between precision and recall is the **F1 score** and gives a unique score. ( $2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$ )
- Other important measures are the True Positive Rate (TPR) and the False Positive Rate (FPR) ( $\text{TPR} = \frac{TP}{TP+FN}$  et  $\text{FPR} = \frac{FP}{FP+TN}$ )
- Read the **Wikipedia page** !!!

# Decision function

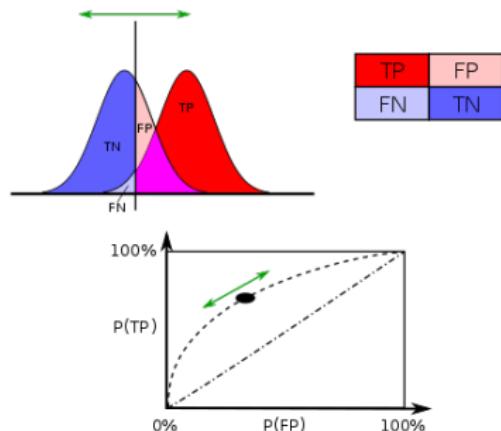
- To obtain predictions, a score is calculated and then the score is compared to a threshold.
- The threshold can be changed and seen as a hyperparameter

# ROC curves

- La courbe ROC (Receiver operating characteristic) trace FPR et TPR dans un diagramme en fonction du seuil.



- When the output of the classifier is considered as a random variable



- The area under the ROC curve can be interpreted as the probability of the classifier giving a higher score to a positive example than to a negative example.

## In the case of multiclass classification

- The measures also work in a multi-class context, with certain restrictions.
- FPR, TPR, etc are calculated per class
- Averages are calculated across classes
- The average is not a good indicator
- In real problems, macro or micro average techniques can be used to correct the bias introduced by the mean
- For example
  - ▶ Class A: 1 TP and 1 FP; Class B: 10 TP and 90 FP; Class C: 1 TP and 1 FP; Class D: 1 TP and 1 FP.
  - ▶ macro-average:  $(0.5 + 0.1 + 0.5 + 0.5)/4 = 0.4$
  - ▶ micro-average:  $(1 + 10 + 1 + 1)/(2 + 100 + 2 + 2) = 0.123$
- We can also weight the averages by the proportions in each class.

# With sklearn

- Functions are listed in `sklearn.metrics`
  - ▶ `confusion_matrix`
  - ▶ `precision_score`, `recall_score` et `f1_score`,
  - ▶ `classification_report` : un tableau avec les principales mesures
- Most of the classifiers have a method `decision_function` or `predict_proba`
- You can draw curves with `precision_recall_curve` and `roc_curve` and the confusion matrix with `plot_confusion_matrix`,  
`plot_precision_recall_curve`, `plot_roc_curve`

# Outline

- 1 Decision Trees
- 2 Evaluation of classification models
- 3 Ensemble methods

# Principle

- (weighted) vote of a set of binders: the majority vote wins the decision
- theoretical results given by **boosting**.
- Intuition: If the basic classifiers are sufficiently independent to be less wrong than a random classifier, then their vote forms a more accurate method than the basic classifiers.
- Numerous methods can be used to combine these classifiers: Bagging, Pasting, Boosting, Stacking, etc.

# Bagging

- You can combine classifiers of a different nature or the same classifiers obtained on different data (but from the same distribution).
- The aim is to reduce variance and limit overfitting
- **bagging**: uniform random selection of a subset of the data (bootstrap aggregating);
- **pasting**: same as above, but drawn without replacement.
- **random subspace**: similar to bootstrap but drawn on attributes and not on examples.
- **random patches**: when you draw on both examples and attributes.
- Aggregation: by mode (highest frequency of predictions), by majority vote, by the average of decision functions.
- Scaling is easier if you can distribute the calculations

# Random Forests

- It's basically bagging with decision trees
- Introducing variety into the construction of a set of decision trees
- The randomness is obtained by
  - ▶ drawing a sub-sample of examples, or
  - ▶ in the search for the best test: randomly mixed attributes, chosen from a subset of the attributes
- The readability of decision trees is lost
- However, the importance of attributes can be calculated as the average gain over all the trees in the forest, and this importance can be visualised.

# Boosting

- modify the distribution of examples by using a weight on each example that varies according to the rate of correct classification
- the weight of poorly classified examples increases gradually
- the decision is a weighted vote. The weight is a function of the accuracy of the classifier

# Boosting principles

- Reminder: error in generalisation for ERM = approximation error + estimation error
- We have a class of poor functions with a large approximation error: **weak** learners whose performance is **just slightly better than random**.
- How can we iteratively make it richer?
- **Boost** we aggregate weak learners to form stronger learners, whose performance can be as large as we want.
- Objectives:
  - ▶ a method that breaks down the complexity of learning a strong learner directly
  - ▶ to answer the theoretical question of how to transform a weak learner into a strong learner.
- **Adaboost** is an implementation that solves this theoretical problem.

# Weak learner

## Definition (Weak learner)

$A$  is a  $\gamma$ -weak-learner for a class  $\mathcal{H}$  if there exists a function  $m_{\mathcal{H}} : (0, 1) \rightarrow \mathbb{N}$  such that for any  $\delta \in (0, 1)$ , for every distribution  $\mathcal{D}$  over  $\mathcal{X}$  and for every labeling function  $f : \mathcal{X} \rightarrow \{-1, 1\}$ , if the realization assumption holds with respect to  $\mathcal{H}$ ,  $\mathcal{D}$ ,  $f$ , then when running the learning algorithm  $A$  on  $m \geq m_{\mathcal{H}}(\delta)$  iid examples generated by  $\mathcal{D}$  and labeled by  $f$ , the algorithm returns a hypothesis  $h$  such that, with probability of at least  $1 - \delta$ ,  $L_{\mathcal{D}, f}(h) \leq 1/2 - \gamma$ .

- $\epsilon$  (strong learner) is replaced by  $1/2 - \gamma$ .

# Adaboost

- we are given a weak learner
- Adaboost is an iterative algorithm that modifies the distribution  $D^{(t)}$  over  $S$  at each round  $t$
- the error of the hypothesis  $h_t$  built by the weak learner is  $\epsilon_t = \sum_i^m D_i^{(t)} \mathbb{1}_{[h_t(x_i) \neq y_i]}$  with probability  $1/2 - \gamma$ .
- The weight associated with  $h_t$  is inversely proportional to the error:  $w_t = \frac{1}{2} \log(\frac{1}{\epsilon} - 1)$  : wrongly classified examples are given a higher weight.
- the output of Adaboost is a weighted vote of all hypotheses  $h_t$

# Algorithm

$$\mathcal{D}^{(1)} = (1/m, \dots, 1/m)$$

**for**  $t = 1, \dots, T$

$$h_t = WL(\mathcal{D}^{(t)}, S)$$

$$\epsilon_t = \sum_i^m D_i^{(t)} \mathbb{1}_{[h_t(x_i) \neq y_i]}$$

$$w_t = \frac{1}{2} \log\left(\frac{1}{\epsilon_t} - 1\right)$$

$$D_i^{(t+1)} = \frac{D_i^{(t)} \exp(-w_t y_i h_t(x_i))}{\sum_j^m D_j^{(t)} \exp(-w_t y_j h_t(x_j))}$$

**return**  $h_s(x) = \text{sign}\left(\sum_{t=1}^T w_t h_t(x)\right)$

- The hypothesis class of Adaboost, given a weak learner that builds hypothesis in  $B$  is  $\{\text{sign}\left(\sum_{t=1}^T w_t h_t(x)\right) \mid \mathbf{w} \in \mathbb{R}^T, h_t \in B\}$
- Theorem : The training error of Adaboost is bounded by  $\exp(-2\gamma^2 T)$

# Gradient boosting

- Another way of focusing on errors: the next classifier in the loop attempts to correctly predict residual errors.
- Example with `DecisionTreeRegressor`.
  - ▶ An instance `dtr` of this classifier produces a residual  $y2 = y - dtr.predict(X)$ .
  - ▶ A `dtr2` is constructed to predict  $y2$ .
  - ▶ Combine the sum of the predictions from the 2 classifiers.
- The `GradientBoostingRegressor` class does this.

# Stacking

- Bias reduction by classifier compositions
- we learn to combine the predictions of basic classifiers
- learning is divided into two parts:
  - ▶ part  $A$  trains  $n$  base classifiers
  - ▶ part  $B$  is used to create a set  $P$  of prediction vectors of dimension  $n$
- the set  $P$  is used to train a combination (*blender*)

# With sklearn

- a VotingClassifier class for simple weighted voting or no voting of classifiers
- the AdaBoostClassifier class
- the BaggingClassifier class
  - ▶ Bagging/pasting setting via the bootstrap parameter
  - ▶ Setting of `max_features` and `max_samples`.
- the RandomForestClassifier class
- the StackingClassifier class
- In these implementations, the aggregation is the average of the decision functions.
- The same classes exist for regression.